

Excitation-Aware On-Track System Identification at Full Scale: A Simulation Study

Ebrahim Abdelghfar Ebrahim¹, Hazem Abuelanin², Mohamed Abdelaziz¹, Mohamed Omar Mohamed Elsayed¹, Sherif Said¹

Abstract—On-track system identification learns tire models from ordinary racing laps instead of dedicated manoeuvres. We study one such residual-learning pipeline at full vehicle scale, in a simulator publishing per-wheel forces as ground truth. Transplanted unchanged from a 1:10 platform it converges on a tire with less than half the vehicle's grip ceiling and most Pacejka coefficients railed. The cause is a loss of excitation created by the pipeline's own virtual-data step, which no change of solver or tuning removes. The friction the synthetic rollout can demand, $\mu_{\text{reach}} = v^2 \delta_{\text{sweep}} / (gL)$, is a property of the rollout configuration and not of the tire, as a grid on a known tire confirms. Two further defects share that source: the fit's objective depends on the coefficients only through the product BCD , so the peak factor is not recoverable from the rollout at any excitation, and the loop feeding each fit its own previous answer is not a contraction. We correct all three (rollout speed from the measured lap, sweep bounded to the observed steering, peak factor measured off the trace and pinned, update damped into a relaxed fixed-point iteration) and gate the controller interface on derived physical quantities rather than on coefficient bounds. Over three timed laps on a surface 33 % below the one its prior was measured on, no coefficient rails, axle cornering stiffness is recovered to 3 % and the peak factor to 11 %, against a railed peak factor and a 125 % stiffness error for the inherited configuration. All results are in simulation, one run per configuration.

Index Terms—Autonomous vehicles, system identification, parameter estimation, vehicle dynamics, tires, predictive control, road vehicles, simulation.

I. INTRODUCTION

MODEL-BASED controllers for autonomous racing derive their performance from the fidelity of the vehicle model they carry [1], [2], [3]. In the lateral channel that fidelity is dominated by the tire, simultaneously the physical bound on manoeuvrability and the least accessible component to measure [4], [5]. Classical characterisation takes the vehicle off the track for steady-state circular or ramp-steer manoeuvres, at a logistical cost a race weekend rarely affords; identifying on track removes the logistics but brings back transients, noise and a racing line that is a poor excitation signal, and nonlinear least squares applied directly to lap data is brittle in exactly this regime [6].

Dikici *et al.* [6] propose the hybrid remedy that is the starting point of this paper: a small residual network trained on 30 s of ordinary lap data predicts the one-step error of a nominal single-track model; the corrected model is rolled

out through a *synthetic* quasi-steady steering sweep; Pacejka coefficients are fitted to that sweep; and each fit becomes the next nominal model. Because the sweep is generated rather than driven it can be made steady and noise-free, and the residual network only has to interpolate within its training distribution. On a 1:10 platform the procedure reports a $3.3 \times$ lower one-step RMSE than nonlinear least squares under noise [6].

A. Problem statement

This paper asks what the pipeline identifies when it is moved from a 3.7 kg, 0.32 m-wheelbase platform to a full-scale Formula Student vehicle: 269.6 kg, 1.53 m wheelbase, 16° of road-wheel lock, following a reference speed profile of 9.6–25.8 m/s. The question is identifiability, not parameter transfer, and two scale-dependent mechanisms, both latent at 1:10, break the estimate. First, excitation collapses inside the virtual-data step: the Pacejka fit never sees the recorded lap, only the synthetic rollout, whose reachable friction $\mu_{\text{reach}} \approx v^2 \delta_{\text{sweep}} / (gL)$ is a property of the rollout configuration and not of the tire, and a rollout speed inherited from the scaled platform caps it at 0.43 here. Below the tire's peak the Magic Formula degenerates to $F_y \approx F_z(BCD)\alpha$, only the product BCD is identifiable, and the bounded optimiser resolves the remaining freedom on a box edge [7], [4]. Second, lateral demand scales as v^2 , so the tracking cost of a bad grip estimate grows quadratically with speed: one identified set with a peak friction coefficient of 0.40 tracked at 0.16 m of cross-track error at 11.1 m/s and 41.89 m at 27 m/s through the same controller (Section VII-A), and the scaled platform never reaches the speed at which that error becomes observable.

One difference works in our favour. Because the study runs in the CARLA simulator [8] through a purpose-built ROS 2 interface, the simulator's own per-wheel telemetry (slip angle, normal load and lateral force) is available as ground truth, so an on-track identifier can be scored against the forces the plant actually generated.

B. Contributions

- 1) **An identifiability analysis of the virtual-data step** (Section V-B): the closed-form reachable-friction bound above, and an *excitation-aware rollout* that selects the rollout speed from the measured lap and warns when the configuration cannot reach the prior's peak. A grid over rollout speed, sweep amplitude and true peak factor on a known tire verifies it (Section VI-B). This corrects the

¹The authors are with the Faculty of Engineering, Ain Shams University, Cairo, Egypt, and the Autotronics Research Lab (ARL), Cairo, Egypt.

Corresponding author: Ebrahim Abdelghfar Ebrahim (e-mail: ebrahimabdelghfar550@gmail.com).

baseline method [6] at any scale, and stays latent at 1:10 because a 2–4 m/s rollout is representative there and is not here.

- 2) **Two structural properties of the loop, and the corrections they imply** (Section V-D): stage 3's residual depends on B , C and D only through their product, so the peak factor is not recoverable from the rollout *at any excitation* and must be measured on the recorded trace and held; and the loop is a positive feedback path rather than a contraction, so stage 4 becomes a relaxed fixed-point iteration at $\beta = 0.20$.
- 3) **The corrections the bound implies**, on both sides of the fit (Sections V-C and V-E): a sweep bounded at $1.5\times$ the 99th percentile of the observed steering, a Tikhonov target frozen at the configured prior, slip angles built from the *achieved* road-wheel angle, mass and wheelbase from measured static wheel loads, a friction warm start from the *curvature* of the axle force curve, and physics-informed dynamics residuals normalised to the state-increment units of the data term, which the published form leaves a factor $2.3e5$ apart.
- 4) **Admissibility gates on derived physical quantities** at the model-to-controller boundary (Section VII): axle cornering stiffness, understeer sign and critical speed, and grip ceiling against the trajectory's lateral demand. They replace per-coefficient box checks, which cannot detect a degenerate fit or a pairwise-oversteering axle set.
- 5) **A simulation architecture that supplies the ground truth, and a closed-loop evaluation on it** (Sections IV-A–IV-D and VI). The evaluation runs on a static 0.70 friction surface, deliberately not the 1.05 the configured tire prior was measured on, so a pipeline unable to move off its prior scores badly. Over three timed laps the corrected pipeline identifies axle cornering stiffness to 3.2 % front and 1.4 % rear and tracks the peak factor to 11 %, against an inherited configuration that rails D on its 2.0 bound in every cycle and overstates stiffness by 118–132 %. Axle force RMSE is $12.1\times$ and $10.5\times$ lower, and MPC lateral tracking error 32 % lower, at unchanged solver cost.

All results are obtained in simulation, and Section VIII states which conclusions are expected to transfer to hardware.

II. RELATED WORK AND POSITION

A. On-track and data-driven identification

The Magic Formula [9], [4] is the model identified here. Classical characterisation obtains its coefficients from steady-state manoeuvres in a controlled space [6], [5]: clean data at high logistical cost. Direct fitting to measured contact-patch forces needs a force-measuring hub or a simulator's own wheel telemetry; here it is used only as a *reference* against which the on-track identifier is scored, never inside the identification loop.

On-track methods trade excitation quality for operational simplicity. Nonlinear least squares on lap data is the classical baseline and is brittle under noise [6]. The recursive alternative

is a dual or joint filter carrying states and parameters together [10], [11], and it is subject to the same excitation condition this paper is about: a filter run on a smooth racing line has as little information about the tire's peak as a batch fit of the same data, and merely expresses that as a parameter covariance that never shrinks rather than as a coefficient on a bound. Dedicated friction estimators make the excitation dependence explicit [12], [13], [14], [15], and Hsu *et al.* [16] is the closest prior art to the warm start of Section V-E2: they recover the friction limit below the peak from the *curvature* of the force response, read through the pneumatic trail in steering torque. Our estimator uses the same observation on the axle force curve reconstructed from the IMU and odometry, with no steering-torque sensor; it is then gated, applied to both axles from the axle the manoeuvre identifies, and pinned for the co-identification loop, which is a property of what that loop can identify (Section V-D1) rather than of the estimator. Gaussian-process approaches learn the mismatch of a nominal model [17], [18], [2] at a continuing computational cost; neural residuals replace the Gaussian process with a cheaper regressor at the cost of generalisation guarantees.

Deep Dynamics [19] is the closest work on the physics-constraint axis: it estimates Pacejka coefficients inside a neural network and adds a *Physics Guard* layer clamping them into nominal ranges. Our admissibility gates share the motivation but differ in placement. We report as a substantive finding that per-coefficient box constraints are insufficient: the degenerate fit that motivated this work satisfied every box constraint while five of eight coefficients sat exactly on a bound, and a later fit satisfied every box constraint while being pairwise oversteering and open-loop unstable above 14.5 m/s. We therefore gate on derived quantities at the model-to-controller interface and treat a coefficient resting on a bound as a diagnostic of non-identifiability rather than a successful clamp. The hybrid residual-plus-virtual-data structure of [6] is the pipeline we extend; the idea it rests on, to query the learned model only where it interpolates and extract physically parameterised coefficients from it, is retained unchanged, and what we correct is the mechanism deciding *where* the learned model is queried.

B. Overlap with concurrent work

Wu *et al.* [20] extend the same baseline concurrently and independently, introducing a vision-based friction prior, an S4 structured state-space residual in place of the MLP, and Nelder–Mead for the iterative Pacejka extraction. The pipeline described here independently contains an S4D temporal residual [21], [22], a friction warm-start, and Nelder–Mead and differential-evolution solver options, so we claim none of the three as contributions. Three differences are technical rather than chronological. Our friction prior is proprioceptive, so it needs no camera and no textured-surface assumption. It is read from the *curvature* of the axle force curve by a gated brush-model fit [23], [14] rather than from utilised friction, which on a gentle lap measures 0.044 median on this vehicle and is only a lower bound (Section V-E2a). And it is *pinned* rather than used as an initialisation, because Section V-D1 shows the

rollout fit cannot separate the peak factor from the cornering stiffness at any excitation, so a prior on D the fit is free to move is overwritten within one cycle. Neither prior work is the fix: the degeneracy is located in the rollout excitation, not the initialisation, and to our knowledge the reachable-friction bound of Section V-B and its consequence for the virtual-data step have not been stated before.

C. Simulation platform and downstream consumer

CARLA [8] supplies the sensor and scenario infrastructure used here. Its native dynamics come from the underlying PhysX wheeled-vehicle model [24], a saturating friction-and-stiffness tire law (Section IV-B) rather than an empirical Magic Formula fit, which is why this study treats the simulator as a well-instrumented *plant* to be identified and not as a source of ground-truth Pacejka coefficients. R-CARLA [25] makes the dynamics model interchangeable, and Sagmeister *et al.* [26] find that a controller's sensitivity to model fidelity is largest near the acceleration limits, where this study operates; CARLA has also hosted complete Formula Student Driverless stacks with ROS 2 interfaces [27]. Scaled platforms, FITENTH [28] and the ForzaETH Race Stack [3], provide the software context of the baseline; full-scale Formula Student Driverless systems [5] the vehicle class targeted here. The identified model is consumed by a nonlinear MPC path-tracking controller built on *acados* [29] with partial-condensing HPIPM [30] [1], [2], tracking a minimum-curvature raceline [31]. We claim no contribution to the controller; it appears only as the consumer that defines what an admissible identification is (Section VII).

III. IDENTIFICATION PROBLEM

A. Model structure

The identification model is the dynamic single-track model of the baseline, with mass m , yaw inertia I_z and centre-of-gravity distances l_f, l_r , $L = l_f + l_r$. Longitudinal and lateral dynamics are decoupled and load transfer is neglected:

$$\dot{v}_y = \frac{1}{m}(F_{y,r} + F_{y,f} \cos \delta - m v_x \omega), \quad (1)$$

$$\dot{\omega} = \frac{1}{I_z}(F_{y,f} l_f \cos \delta - F_{y,r} l_r), \quad (2)$$

with slip angles

$$\alpha_f = \delta - \arctan\left(\frac{v_y + l_f \omega}{v_x}\right), \quad \alpha_r = -\arctan\left(\frac{v_y - l_r \omega}{v_x}\right). \quad (3)$$

Axle forces follow the lateral Magic Formula

$$F_y = F_z D \sin(C \arctan(B\alpha - E(B\alpha - \arctan B\alpha))), \quad (4)$$

with per-axle coefficients $\Phi_p = [B_f, C_f, D_f, E_f, B_r, C_r, D_r, E_r]$. The identification state and input are $\mathbf{x} = [v_y, \omega]$ and $\mathbf{u} = [v_x, \delta]$, and the nominal one-step map is $\hat{\mathbf{x}}_{k+1} = f(\mathbf{x}_k, \mathbf{u}_k, \Phi_p)$. The baseline discretises (1)–(2) by explicit Euler at T_s ; we keep Euler as the default and expose Runge–Kutta as an option, because the full-scale vehicle's lateral mode is stiffer relative to T_s than the scaled platform's.

B. Objective and derived quantities

The objective is to recover Φ_p from a single 30 s segment of ordinary lap data, without dedicated manoeuvres, such that the resulting model is both predictive and *admissible* for the controller that consumes it. Admissibility is expressed on derived quantities rather than on coefficients. The linearised axle cornering stiffness is the slope of (4) at zero slip,

$$C_i = F_{z,i} B_i C_i^{\text{MF}} D_i, \quad i \in \{f, r\}, \quad (5)$$

(writing C^{MF} for the shape factor to avoid a clash), and the grip ceiling implied by an identified set is

$$a_y^{\text{max}} = \frac{D_f F_{z,f} + D_r F_{z,r}}{m}. \quad (6)$$

The linear model built from (5) is oversteering when $l_f C_f > l_r C_r$ and is then open-loop unstable above the critical speed $v_{\text{crit}} = L \sqrt{C_f C_r / (m(l_f C_f - l_r C_r))}$ [32].

C. Practical identifiability under limited excitation

Two properties of (4) are used throughout. First, its slope at zero slip is $F_z BCD$, so the cornering stiffness depends on the coefficients only through that product, and E drops out. Second, if the data never approaches the peak, (4) linearises to $F_y \approx F_z(BCD)\alpha$ and B, C, D become mutually unidentifiable. Given data that does reach the peak, (4) is *structurally* identifiable [33]; what fails here is *practical* identifiability [34], the standard persistent-excitation condition [7], [35] written on this model, and we claim neither result as novel. Our claim, demonstrated in Section V-B, is that the baseline pipeline's virtual-data step can silently violate that condition, and that a bounded optimiser then resolves the remaining freedom by sliding onto a box edge. The violation surfaces as a converged fit rather than a failure.

D. Baseline pipeline

The baseline repeats four stages for N_{it} iterations [6]. A small network is trained on the recorded lap to predict the nominal model's one-step error from $[v_x, v_y, \omega, \delta]_k$ (stage 1); the corrected model is integrated open-loop at constant speed while the steering angle is ramped from 0 to δ_{sweep} , producing a synthetic quasi-steady dataset (stage 2); assuming $\dot{v}_y = \dot{\omega} = 0$ on that data, axle forces are recovered from the yaw balance,

$$F_{y,r} = \frac{m l_f v_x \omega}{L}, \quad F_{y,f} = \frac{m l_r v_x \omega}{L \cos \delta}, \quad (7)$$

and Φ_p is fitted to $(\alpha_i, F_{y,i})$ by bounded least squares (stage 3); the fitted Φ_p then becomes the next iteration's nominal model (stage 4). Algorithm 1 gives the corrected cycle in full.

The rest of this paper follows from one property of that structure: stage 3 never sees the recorded lap, only the output of stage 2. The recorded lap enters only through the residual network, so the excitation available to the fit is set by the rollout configuration and not by how the vehicle was driven.

TABLE I
VEHICLE CONFIGURATION AND MEASURED PLANT PROPERTIES.

Quantity	Value	Unit
Body mass (configured)	240.0	kg
Static wheel load (total)	2644.5	N
Total (kerb) mass m^a	269.6	kg
Front/rear static split	51.2 / 48.8	%
Wheelbase L	1.5334	m
l_f / l_r	0.7381 / 0.7953	m
Yaw inertia I_z^c	51.1	kg m ²
Max road-wheel angle	16.0	°
Configured wheel friction ^d	1.0	—
Road-surface factor	0.70	—
Effective friction ^e	0.70	—
Realised peak axle friction	0.71	—

^aTotal static wheel load divided by g ; this is the mass every equation in this paper uses, not the configured body mass. ^cAssumed; see Section A. ^dA scenario parameter, not a vehicle property. ^eConfigured wheel friction multiplied by the road surface's own coefficient, as reported per wheel. It is the effective value, not the configured one, that the identifier is scored against (Section V-E2a).

IV. SIMULATION ARCHITECTURE AND DATA PIPELINE

The pipeline is exercised end to end inside a single architecture: a Formula Student vehicle model in CARLA acting as the plant, a C++ ROS 2 communication layer that drives it and republishes its state, and an identification stack that consumes that state. CARLA [8] was selected because it exposes per-wheel slip angle, normal load and tire force rather than chassis states alone, and because its synchronous mode with a fixed physics step (0.033 s here) stamps every sample on the simulation clock; those two properties are what make the identifiability claims of Section V checkable rather than arguable [36].

A. Plant: Formula AI vehicle in CARLA

The plant is the driverless Formula Student-class vehicle model, instantiated in CARLA (server 0.9.16, Unreal Engine 4.26) on the *silverstone* map: a four-wheeled, front-steered, all-wheel-drive rigid-body vehicle whose lateral forces are generated by the PhysX wheeled-vehicle tire model [24] under a configured friction coefficient and lateral stiffness, not by a Magic Formula. The simulator therefore has no “true” Pacejka coefficients to recover; instead the lateral forces PhysX computes at each wheel are taken as ground truth, and the identified Magic Formula is validated against those forces.

B. The plant's tire curve, in closed form

The identification target is known here in closed form. PhysX computes the lateral force of each wheel with `PxVehicleComputeTireForceDefault` [24]; at zero longitudinal slip and zero camber that expression collapses to

$$F_y(\alpha) = \text{sgn}(\alpha) \mu F_z S_1(K), \quad K = \frac{C_{\text{lat}} |\tan \alpha|}{\mu F_z}, \quad (8)$$

with $C_{\text{lat}} = F_z^{\text{rest}} k_y S_1(3\bar{F}_z/k_x)$, $\bar{F}_z = F_z/F_z^{\text{rest}}$ the normalised load, (k_y, k_x) the wheel's lateral-stiffness parameters, and $S_1(K) = \min(1, K - K^2/3 + K^3/27)$ the engine's smoothing function.

TABLE II
THE PLANT'S LATERAL TIRE CURVE: THE PARAMETERS OF (8), READ FROM THE SIMULATOR'S OWN WHEEL CONFIGURATION AND TELEMETRY.

Quantity	Value	Unit
Effective friction μ	0.70	—
Lateral stiffness k_y	17.0	/rad
Max normalised load k_x	2.0	—
Initial slope C_{lat} at $\bar{F}_z = 1$	14.875 F_z^{rest}	N/rad
front axle, $F_{z,f} = 1373$ N	20.4	kN/rad
rear axle, $F_{z,r} = 1274$ N	19.0	kN/rad
Peak reached at, $\bar{F}_z = 1, \mu = 0.70$	8.0	°

The initial slope is set by the *rest* load and by the normalised load through S_1 , not by F_z alone, and falls monotonically with $\bar{F}_z$. Every axle slope quoted in this paper is taken at $\bar{F}_z = 1$ and at the static axle loads listed here.

The peak angle follows from $K = 3$ in (8) and therefore moves with μ : 8.0° at the effective $\mu = 0.70$ of the runs reported here, 12.0° at the $\mu = 1.05$ of the prior-fitting run.

Two properties of (8) matter downstream. It saturates at μF_z , first reached at $K = 3$, and it never falls off past that peak, so a Magic Formula can match it up to the peak and no further: E and the post-peak decay exist to describe a falling branch this curve does not have. Its scale is set entirely by μ and C_{lat} , both of which the simulator publishes, so the curve of Table II is read off the plant rather than fitted to it.

That the closed form is the right one was checked rather than assumed: on 1804 ground-truth wheel samples from a live lap ($P_{99}(|\alpha|) = 3.66^\circ$), each predicted from its own measured α , F_z and μ , (8) returns $R^2 = 0.99$ at 16.1 N front and 14.7 N rear RMSE. Section VI-D uses this curve to separate the error of the identified fit from the error of representing (8) in the form (4) that the controller requires.

C. CARLA_ASU_BRIDGE: interfaces and telemetry

CARLA_ASU_BRIDGE [37] is a C++ ROS 2 lifecycle node linked directly against CARLA's client library (Fig. 1). It spawns and configures the ego vehicle, drives its inputs over a single `ackermann_msgs/AckermannDriveStamped` on `/drive`, and republishes simulator state as standard ROS 2 topics; all identification traffic crosses this interface. On the feedback path, one property is specific to this architecture and the rest of this paper rests on it: the simulator's per-wheel slip angle, normal load and lateral force are republished on the same simulation clock as the odometry the identifier consumes, so an identified tire model can be scored against the forces the plant actually generated rather than a second model. Appendix B lists the full interface and the body-frame and sign conventions at the simulator-ROS 2 boundary.

D. Data acquisition and preprocessing

A model-free pure-pursuit controller drives the vehicle around the reference raceline; it requires no tire knowledge beyond the wheelbase and is therefore a valid bootstrap, as in the baseline [6], [3]. Odometry, the achieved steering angle and the drive command are sampled into a ring buffer by a 50 Hz timer for 30 s, gated on $v_x > 1$ m/s. The timer runs on the simulation clock, which advances only at the server's 0.033 s step, so the achieved acquisition rate is the step rate.




Fig. 1. Architecture of the CARLA_ASU_BRIDGE ROS 2 interface. The bridge links against CARLA's C++ client, drives the ego vehicle through CARLA's native Ackermann controller, and republishes state, sensors and per-wheel telemetry on ROS 2 topics stamped from the simulation clock. The *feedback/tire_forces* path (highlighted) carries the simulator's own per-wheel slip angle, normal load and lateral force, and is the ground truth against which the identifier is scored.

Two mismatches follow. The acquisition timer asks for a sample every 0.020 s, 65 % more often than the server produces one, so a fraction of consecutive buffer entries are repeated states and the one-step residual target carries a zero-order-hold artefact on those pairs. Separately, the one-step model of Section III-A is discretised at $T_s = 0.035$ s, 6 % longer than the physics step. T_s is empirical and is carried because the identification proved insensitive to it: at 0.020 s and 0.035 s the corrected configuration returns the *same* last accepted coefficient set to four decimal places, and the inherited configuration keeps D on its 2.0 bound at both, so the box-edge signature is not an artefact of the mismatch. Results are reported at $T_s = 0.035$ s. Preprocessing follows the baseline: a zero-phase low-pass filter, then symmetry augmentation by mirroring $(v_y, \omega, \delta) \rightarrow (-v_y, -\omega, -\delta)$ at unchanged v_x . Re-identification retriggers every 30 s by default.

Only the deployment stage leaves this architecture: validation draws on *feedback/tire_forces* and *feedback/imu*, neither of which enters the identification loop, so the estimator is scored against a channel it did not observe, and Section VI can attribute a closed-loop difference to the identified model rather than to the interface carrying it.

V. IDENTIFICATION METHODOLOGY

Figure 2 summarises the loop and Algorithm 1 states it in the form needed to reimplement it; lines marked $\triangleright$ are this paper's changes to the inherited pipeline. Sections V-B–V-C carry the central methodological contribution.

A. Residual model and objective

The residual carries whatever the nominal model (1)–(4) does not represent. The default is the baseline multilayer

```

in: buffer  $D = \{v_x, v_y, \omega, \delta^{\text{ach}}\}_{1:K}$  (30 s), prior  $\Phi_p^0$ , bounds  $[\Phi_p^-, \Phi_p^+]$ 
out:  $\Phi_p$ , or UNIDENTIFIED
1:  $\delta \leftarrow$  achieved road-wheel angle, command on staleness  $\triangleright$ 
2:  $m, l_f, l_r \leftarrow$  measured static wheel loads  $\triangleright$ 
3:  $\hat{\mu}_f \leftarrow \text{BRUSHFIT}(D)$ , front axle, released only if gated  $\triangleright$ 
4: if released then  $D_f^0, D_r^0 \leftarrow \hat{\mu}_f$ ;  $D$  held for the whole loop  $\triangleright$ 
5:  $\Phi_p^{\text{tik}} \leftarrow \Phi_p^0$  (frozen for all iterations)  $\triangleright$ 
6:  $v \leftarrow \text{clip}(\bar{v}_x(D), v_{\min}, \infty)$   $\triangleright$ 
7:  $\delta_{\text{sweep}} \leftarrow \text{clip}(1.5 P_{99}(|\delta|), \delta_{\min}, \delta_{\max})$   $\triangleright$ 
8:  $\mu_{\text{reach}} \leftarrow v^2 \delta_{\text{sweep}} / (gL)$ ; warn if  $\mu_{\text{reach}} < \max_a D_a^0$   $\triangleright$ 
9:  $\Phi_p \leftarrow \Phi_p^0$ ; for  $i = 1$  to  $N_{\text{it}}$  do
10:  $f_{NN} \leftarrow \text{TRAINRESIDUAL}(D, \Phi_p)$ ;  $S \leftarrow$ 
    ROLLOUT( $\Phi_p, f_{NN}, v, \delta_{\text{sweep}}$ )
11:  $\Phi_p \leftarrow \arg \min_{B,C,E} \|F_y(\Phi_p; S) - \hat{F}_y(S)\|^2 + \lambda \|\Phi_p - \Phi_p^{\text{tik}}\|^2$ ,
     $D$  held  $\triangleright$ 
12:  $\Phi_p \leftarrow (1 - \beta) \Phi_p + \beta \tilde{\Phi}_p$  ( $\beta = 0.20$ )  $\triangleright$ 
13: if any  $\Phi_p$  component on a bound then report the degeneracy  $\triangleright$ 
14: if not GATES( $\Phi_p$ ) then return UNIDENTIFIED  $\triangleright$ 
15: submit  $\Phi_p$ ; keep the previous model until acknowledged  $\triangleright$ 
: GATES: axle stiffness sign and magnitude;  $l_f C_f < l_r C_r$  and  $v_{\text{crit}}$ 
outside the speed range; grip ceiling  $\geq$  the trajectory's lateral demand
(Section VII).

```

Algorithm 1: Corrected on-track identification cycle. $\triangleright$ marks a change to the inherited pipeline [6].

perceptron: 4 inputs, one hidden layer of 8 units, 2 outputs, LeakyReLU, 58 parameters, full-batch Adam [38] at 5×10^{-4} on an MSE objective, in PyTorch [39]. The architecture is selectable; the two arms of Section VI use baseline and s_4 , the latter an S4D diagonal state-space residual [21] over a sliding window of $W = 20$ samples (0.7 s), so the residual can carry tire-relaxation and yaw-damping dynamics a per-sample map cannot. We claim neither that architecture, which is also in concurrent work [20], nor the $8.7\times$ reduction in its training cost, an exact algebraic rewrite that changes no identified coefficient; both are in the supplementary material, and the identifiability result below holds for either architecture.

1.5402/1.8721 over six iterations, a claimed 1.7 g on a vehicle measuring 1.0 g in steady state.

Two changes close the loop. First, an extrapolation bound: the sweep terminates at $\delta_{\text{sweep}} = \text{clip}(\lambda_e P_{99}(|\delta_{\text{obs}}|), \delta_{\text{min}}, \delta_{\text{max}})$ with $\lambda_e = 1.5$, $\delta_{\text{min}} = 0.05 \text{ rad}$ and $\delta_{\text{max}} = 0.54 \text{ rad}$, so the network is never queried more than 50 % beyond the steering range it was trained on. Both δ_{max} and the inherited 0.4 rad lie above the plant's 0.279 rad mechanical lock deliberately: the sweep is a *virtual* rollout, and clamping it to the lock would cap μ_{reach} a second time, by the actuator instead of the rollout speed (Section A, A6). Second, a frozen prior: the Tikhonov target is held at the configured Φ_p^0 for all iterations while the optimiser's *start point* continues to track the previous iteration, giving the per-axle objective

$$\min_{\phi} \sum_k \rho(r_k(\phi)) + \sum_j w_j (\phi_j - \phi_j^0)^2, \quad (12)$$

with r_k the force residual, ρ a soft- ℓ_1 robust loss and w_j scaled to the data residual budget by a single dimensionless weight $\lambda = 0.05$. The regulariser exists to hold the coefficients the manoeuvre does not excite, E above all, at their nominal value instead of on a box edge.

D. What the rollout fit identifies, and why the loop must be damped

1) *The rollout fit sees only the product BCD*: In the virtual-data step the low-slip degeneracy of Section III-C is stronger, and structural rather than excitation-dependent. Stage 3 does not observe a force; it reconstructs one from the rollout's own quasi-steady yaw balance (7), which on the synthetic trajectory reduces to

$$\frac{F_{y,f}}{F_{z,f}} = \frac{v_x \omega}{g \cos \delta}, \quad \frac{F_{y,r}}{F_{z,r}} = \frac{v_x \omega}{g}, \quad (13)$$

so the only quantity (12) ever compares against is the lateral acceleration the rollout realises, which the *previous* coefficient set produced through its slope $F_z BCD$. The fit's residual therefore depends on (B, C, D) only through their product, and D is the free coordinate of the hyperbola $BCD = \text{const}$. Replayed on five identification buffers of a live $\mu = 1.05$ run, the product is recovered tightly, $BCDF_{z,f}$ over 18.9–21.1 kN/rad against a plant measured at 20.4 kN/rad, while the D reaching that same product ranges over 0.69–2.00. D must therefore not be asked of this fit at all. It is measured on the *recorded* trace by the brush estimator of Section V-E2, which fits $F_y(\alpha)$ directly and is not subject to (13), then held fixed while B , C and E carry the stiffness the rollout does determine.

2) *The loop is not a contraction*: Stage 4 makes each iteration's answer the nominal model stage 2 rolls out, and stage 3 then fits the trajectory that model produced, over a steering range already extrapolating $1.5\times$ past the observed one. The fit reads its own forces back, so the loop is a positive feedback path on its own extrapolation error rather than a contraction: replayed offline on five recorded buffers, axle force RMSE against the plant's telemetry degrades monotonically with iteration count, 38.9, 53.0, 58.5 and 70.4 N at the front

for $N_{\text{it}} \in \{1, 2, 3, 6\}$. Pinning D does not remove this; it relocates it into B , C and E . The remedy is to damp the step, not to shorten the loop, so stage 4 becomes a relaxed fixed-point iteration

$$\Phi_p^{(i)} \leftarrow (1 - \beta) \Phi_p^{(i-1)} + \beta \tilde{\Phi}_p^{(i)}, \quad (14)$$

with $\beta = 0.20$; $\beta = 1$ recovers the inherited update exactly. The sweep over β (Section VI-I) carries the control the contribution needs: $\beta = 0$ freezes the coefficients at the pinned prior, performing no identification at all, and is the worst setting. The iterations do carry information, but only a damped update preserves it.

Fitting itself applies (7) to the synthetic data under box constraints $\Phi_p \in ([4.0, 1.2, 0.4, -3.0], [20.0, 2.2, 2.0, 1.0])$ per axle, over B , C and E with D held. Relative to the baseline we raised the jittered multi-starts from 1 to 8 and added a derivative-free simplex [41] and a global differential-evolution solver [42] alongside trust-region reflective [43], through the same library [44]. Neither change addresses the degeneracy, and Section VI-B verifies both as secondary.

E. Measurement-side corrections

1) *Achieved steering, measured mass and geometry*: α_f in (3) was built from the `/drive` setpoint. The measured achieved/commanded ratio is 0.76, so α_f was systematically $\approx 25\%$ too large and the fitted front stiffness $\approx 20\%$ too small (13725 N/rad against 17198 N/rad fitted to the simulator's own per-wheel forces over the same run); using `feedback/steering_angle` recovers 18486 N/rad. Separately, m and l_{wb} come from the measured static wheel loads (Table I) rather than a configuration file, and $l_{wb} \equiv l_f + l_r$ is enforced by a regression test, because (7) uses the wheelbase for the normal loads and the axle split for the force division, and the two disagreeing is a silent bias.

2) *Proprioceptive friction warm-start*: The inherited pipeline had no prior for the peak factor: D started from a configuration constant and was left to the fit. Newton–Euler on the single-track body of (1)–(2), solved for the axle forces rather than the accelerations, gives

$$F_{y,f} = \frac{I_z \dot{\omega} + l_r m a_y}{L \cos \delta}, \quad F_{y,r} = \frac{l_f m a_y - I_z \dot{\omega}}{L}, \quad (15)$$

which needs the IMU's lateral acceleration and the yaw acceleration and no tire model; vertical loads carry the longitudinal transfer $F_{z,f/r} = m(g l_{r/f} \mp a_x h_{cg})/L$ [32], and yaw acceleration is the derivative of a local polynomial fit to the yaw rate (Savitzky–Golay, window 21, order 2) [45] rather than a difference, because $I_z \dot{\omega}$ is a $\approx 10\%$ term in $F_{y,f}$ here.

Per axle, cornering stiffness and peak friction are then fitted jointly to the Fiala brush model [23], [46], [4]

$$F_y = \mu F_z (1 - (1 - \xi)^3) \text{sgn } \alpha, \quad \xi = \min\left(\frac{C_\alpha |\tan \alpha|}{3\mu F_z}, 1\right), \quad (16)$$

by bounded nonlinear least squares with a soft- ℓ_1 loss [44], [43]. The brush model replaces (4) here because it has two parameters, both of which the manoeuvre can support, against the four of the model whose non-identifiability at low slip

this paper is about. What beats the utilisation bound is that μ enters (16) through the *curvature* of $F_y(\alpha)$ and not only through its asymptote: writing $r = |F_y|/(C_\alpha|\tan \alpha|)$ for the force's departure from its own linear extrapolation, (16) inverts to

$$\mu = \frac{2k}{3 - \sqrt{3(4r - 1)}}, \quad k = \frac{C_\alpha|\tan \alpha|}{3F_z}, \quad (17)$$

so a 30% departure from linearity ($r = 0.7$) already pins μ , well below the limit.

Equation (17) also states its own failure mode: $d\mu/dr$ grows without bound as the curve stays linear. The estimate is released only when four conditions hold: absolute grip demand $\max |F_y|/F_z \geq 0.40$, peak friction utilisation $\max |F_y|/(\hat{\mu}F_z) \geq 0.40$, relative precision $\sigma_{\hat{\mu}}/\hat{\mu} \leq 0.15$ from the Gauss–Newton covariance, and $\hat{\mu}$ not resting on its own search bound; otherwise the module logs which condition failed and falls back to the utilisation bound. The first two are the same quantity with and without the fit in the denominator, and both are needed: the utilisation test divides by the *fitted* $\hat{\mu}$ of the same least-squares call, so an axle whose μ comes out too low clears the gate on its own error, while the absolute ratio has no such freedom. The 40% floor is the level the friction-estimation literature converges on [14].

a) *Applied two-sided, then pinned:* The released $\hat{\mu}$ replaces D on both axles, in either direction, and D is then held fixed for the whole co-identification loop. Both axles take the *front* fit: the rear slip angle does not carry δ , so on this trajectory the rear force curve has too little curvature to support (17) and its μ comes out high (1.495/1.383/1.257 rear against 1.011/0.964/0.837 front on three admitted buffers of a $\mu = 1.05$ run). This is not a yaw-inertia error: in quasi-steady cornering $F_{y,f}/F_{y,r} \rightarrow l_r/l_f$, the same ratio as $F_{z,f}/F_{z,r}$, so $\mu_f = \mu_r$ whatever I_z is, and there is one road surface. It is two-sided rather than a floor, the rule it replaces having protected against a D that *moved* between cycles rather than one that was low; it is pinned and not merely seeded, since (13) leaves D free along $BCD = \text{const}$ and on a live run a 0.785 seed ended at the 2.000 bound within one cycle; and an axle with no released $\hat{\mu}$ holds the D it came in with, since there is nothing in the rollout to fit it against.

Pinning puts the whole weight of the result on the gate, and σ_{μ}/μ from the fit covariance is a *precision* figure, not an accuracy one: a wrong mass, yaw inertia, IMU bias or road bank all fit tightly and wrongly through (15). Here the configuration file's 240.0 kg against the 269.6 kg the wheels actually carry biased μ 11% low until the measured value replaced it, and the configured wheel coefficient is not the effective one the physics step uses, the server multiplying it by the road surface's 0.70.

F. Deployment and validation protocol

The identifier runs as a single ROS 2 node cycling a rolling FIFO buffer through four states (collect, train, await the controller's acknowledgement, run), re-identifying every $T_{\text{reid}} = 30$ s. Coefficients are submitted over a service call whose non-acknowledgement leaves the previous model in place, so a rejected identification degrades to the last accepted

model rather than to none. That acknowledgement also arms the scoring: a submission the controller *rejects* on its admissibility gate (Section VII) is a model the car never drove on, and the seconds between a fit completing and its acknowledgement landing are seconds still driven on the previous model.

Three levels of validation are used, in increasing strength, and all three score the *accepted* model only. One-step prediction on a held-out run, as in the baseline: RMSE and R^2 of v_y and ω , with online statistics in Welford's form [47] and samples admitted only above 1 m/s. Direct force comparison against `feedback/tire_forces`: the identified curve is evaluated at the measured (α, F_z) and compared per axle against the measured F_y , and against the plant's own curve (8) at the same points, which separates the error of the fit from the error of representing a PhysX tire. Closed-loop tracking under the MPC: lateral deviation, its time-weighted integral, heading error and solver cost, split by active controller so the pure-pursuit acquisition phase is not scored against a model it does not use.

VI. EXPERIMENTS AND RESULTS

A. Setup


All experiments run on the architecture of Section IV: the Formula AI vehicle in CARLA driven through ROS 2 Humble, the server in synchronous mode at a 0.033 s fixed physics step. The identifier collects 30 s of data at 50 Hz and re-identifies from the newest buffer, matching the baseline's data budget [6]. Ground truth is the simulator's own per-wheel telemetry, which the identifier never reads; every figure is exported with the CSV holding its samples, and the cross-run comparison reads those CSVs rather than re-deriving any metric.

1) *Surface:* Every run uses a *static* surface: the per-wheel coefficient is held at its configured 1.0, and the road-surface factor of 0.70 PhysX applies on top of it makes the effective coefficient 0.70, constant to the last digit the simulator reports over all 10744 and 12025 force-validation samples of the two runs. It is deliberately *not* the $\mu = 1.05$ surface the tire prior was measured on (Section VI-A4), so an arm that cannot move off its prior cannot score well here. It remains a *single* surface: the adaptivity claim a friction sweep would support is not made (Section VIII-B).


2) *Reference trajectory:* Both runs follow the same optimised raceline (Fig. 3), generated with the TUM global race trajectory optimisation package [48] under its minimum-curvature formulation [31]: a closed 5828.2 m lap of 2916 waypoints, reference speed 9.62–25.78 m/s, minimum corner radius 18.7 m ($|\kappa|_{\text{max}} = 0.0536/\text{m}$). Its peak lateral demand is $\max(v_x^2|\kappa|) = 4.98 \text{ m/s}^2$ (0.51 g), roughly 73% of the 0.70 g this surface sustains. It replaces an earlier trajectory whose 1.22 g demand the admissibility gate of Section VII rejected against the same plant; that regeneration, not a relaxed gate, unblocked the closed-loop runs.

3) *The two configurations:* Two configurations of the same pipeline are run back to back on that trajectory, three timed laps each, with an untimed out-lap. They differ in the five settings of Table III and in nothing else; vehicle, map, trajectory,

(a) raceline, 5.83 km lap, coloured by lateral demand



(b) reference speed profile



(c) lateral demand against available grip




Fig. 3. The optimised raceline used for both runs [48], [31]. (a) Track layout, coloured by lateral demand, with the marked corner ($R = 18.7$ m) the tightest on the lap. (b) Reference speed profile. (c) Lateral demand $v_x^2/|\kappa|$ along the lap against the flat 5.0 m/s^2 envelope handed to the optimiser (dashed) and the grip the vehicle has on this surface (dotted, $\approx 0.70 \text{ g}$, measured). No segment exceeds the plant's grip.

TABLE III
THE TWO CONFIGURATIONS.

Setting	Baseline-NN-MSE	Excitation-S4D-Pinn
Residual architecture	MLP (baseline)	S4D (s4)
Residual objective	One-step MSE	Physics-informed, (9)
Pacejka solver	Trust-region [43]	Differential evolution [42]
Friction warm start	Off	On; D pinned per axle
Update relaxation β	1.00 (undamped)	0.20

Everything not listed is identical between the two, β is the relaxation of (14); 1.00 is the inherited undamped update.

controller, data budget, loop count and box constraints are shared. Baseline-NN-MSE is the inherited configuration evaluated at full scale on the corrected rollout, so the comparison isolates those five settings rather than reproducing the published pipeline in full; Excitation-S4D-Pinn is the shipped configuration. Since the five settings move together, what follows is a configuration comparison, not a per-factor ablation, which Section VI-I supplies offline. Nor is it a test of this paper's central contribution: the excitation-aware rollout is active in *both* arms, and the correction itself is evidenced in Section VI-B.

4) *Provenance of the tire prior*: The prior Φ_p^0 both arms start from, and that the Tikhonov term of (12) pulls toward, is the measured tire of Table II, obtained by fitting (4) to CARLA's own per-wheel telemetry over a dedicated limit-excitation run that reached $\alpha = 0.37 \text{ rad}$, axle $\mu = 1.05$ and 1.50 g . It is not independent of the validation data, since that is the same channel the force and friction results are scored against; the manoeuvre that produced it is one the pipeline is not allowed to run, the raceline peaking at 0.51 g ; and, decisively, it was measured on a surface these runs do not use, carrying $D = 1.009/1.002$ and axle slopes of 13.2 kN/rad front and 13.9 kN/rad rear against the 0.70 and the 17.4 and 16.6 kN/rad this surface realises. Returning it unchanged is a 44% grip overstatement and a $16\text{--}24\%$ stiffness deficit, so an arm has to move away from that prior, in the right direction,

to score well.

B. Identifiability of the peak factor against rollout excitation

The claim of Section V-B concerns the rollout configuration alone, so it is tested where the answer is known exactly rather than in CARLA. A tire with a known peak factor D is rolled out through the synthetic sweep of stage 2 (the shipped rollout, residual network set to zero) and fitted back with the shipped solver, *unregularised* and started from the true coefficients, so everything it loses is lost to excitation. Rollout speed v , sweep amplitude δ_{sweep} and true D are swept, each cell repeated over five noise realisations at the one-step persistence RMSE measured on CARLA runs of this vehicle. The configuration axes are swept one at a time, because μ_{reach} is a function of both; the repetition is over noise rather than the solver because the noise-free fit is exactly seed-invariant here.

The grid says three things. First, the inherited 4 m/s rollout is uninformative about D at *every* grip level tested: the 0.70 , 1.05 , 1.40 and 1.80 tires all come back as $\hat{D}$ between 0.50 and 0.59 , with two to seven of the eight coefficients on a bound, so the fit reports μ_{reach} and not the tire. Second, the boundary is $\mu_{\text{reach}} \gtrsim D$ and not a speed: at 6 m/s , where $\mu_{\text{reach}} = 0.96$, the 0.70 and 1.05 tires are recovered to within 3% while the 1.40 and 1.80 tires are 13% and 30% short; above $\mu_{\text{reach}} = 1.7$ every tire is recovered to within 4% with no systematic raiing. The bound predicts where the information runs out, not a step. Third, the failure is not the solver's: multi-start, a global solver or a better start point cannot recover a coefficient the data does not contain.

Panel (b) sweeps the other argument of (11) at fixed $v = 15 \text{ m/s}$. Starving the amplitude destroys D just as starving the speed does, so speed is not the governing variable, and the direction of the error is set by the rollout rather than the tire: an under-speeded rollout returns *too little* grip, an under-swept one too much (the 1.05 tire at 1.30 , the 1.80 tire railed on the 2.0 bound at $\delta_{\text{sweep}} = 0.06 \text{ rad}$). That is the signature of

TABLE IV
PEAK FACTOR $\hat{D}$ RECOVERED FROM A KNOWN TIRE, ALONG BOTH ARGUMENTS OF μ_{reach} .

swept	μ_{reach}	$\hat{D}$ (railed of 8) at true D of			
value		0.70	1.05	1.40	1.80
<i>(a) rollout speed v [m/s] swept at $\delta_{\text{sweep}} = 0.40$ rad</i>					
4	0.43	0.52±0.036 (1.8)	0.57±0.013 (2.8)	0.59±0.018 (3.8)	0.50±0.001 (7)
6	0.96	0.68±0.001 (1)	1.04±0.014 (1.2)	1.22±0.066 (1.8)	1.26±0.097 (2)
8	1.70	0.69±0.001 (1)	1.02±0.001 (0.8)	1.36±0.002 (0)	1.79±0.025 (1.2)
13	4.49	0.72±0.001 (0)	1.05±0.002 (0)	1.38±0.001 (0.4)	1.77±0.002 (0.4)
20	10.64	0.92±0.004 (1)	1.14±0.005 (0.4)	1.43±0.002 (0)	1.80±0.005 (0)
<i>(b) δ_{sweep} [rad] swept at $v = 15$ m/s</i>					
0.03	0.45	0.80±0.019 (2)	1.10±0.041 (2)	1.08±0.182 (2)	1.09±0.397 (2)
0.06	0.90	0.92±0.003 (2)	1.30±0.008 (2)	1.68±0.017 (2)	1.80±0.232 (3.4)
0.15	2.24	0.99±0.004 (2)	1.38±0.003 (2)	1.71±0.019 (2)	1.98±0.009 (1)
0.40	5.98	0.76±0.009 (1)	1.05±0.002 (0)	1.39±0.002 (0)	1.78±0.001 (0)
0.54	8.08	0.83±0.004 (1)	1.08±0.002 (0)	1.38±0.001 (0)	1.78±0.008 (0)

Unregularised, started from the true coefficients. Each cell is $\hat{D} \pm$ the half-width of its 95% t -interval over five noise realisations, with the mean number of the eight coefficients resting on a box constraint in parentheses. The two axes of one rollout are not independent, so each realisation contributes the mean of its own ($\hat{D}_f$, $\hat{D}_r$) and $n = 5$.

μ_{reach} is (11) at $L = 1.5334$ m; the inherited configuration is the first row of (a). It is the Ackermann upper bound and above $\mu_{\text{reach}} \approx 2$ detaches from what the rollout realises, so a large value means “far past saturation”, not a demand the tire ever sees.

an unidentified parameter, not a biased one. The axes do *not* collapse onto a common μ_{reach} , for the upper-bound reason of Section V-B: at $v = 20$ m/s (11) reads 10.64 against a realised 1.40–2.18. Read against the realised demand both axes behave alike, so (11) should be used as the pipeline uses it: a cheap screening bound, conclusive when it falls below D , and not a measure of excitation achieved. One property bounds how all of this should be read: with the residual set to zero, a Magic Formula is fitted to Magic-Formula data under exactly matched structure, so the recovery is a best case.

C. Cross-run summary

Table V is the complete comparison. Excitation-S4D-Pinn is better on every quantity scored against the plant (peak friction, axle force on both axles, both one-step states, lateral tracking error in every phase, and both error tails) and behind only on RMS QP solve time, by 5% of a figure two orders of magnitude inside the control period. The margins are 9.0–12.1× on peak friction and axle force and 1.5× on MPC tracking error. The rest of this section reads those rows against the figures; Section VI-I attributes them.

D. Identified tire model

Figure 4 puts both identified curves against the plant’s own curve (8) and lists the coefficients. Three things are visible.

First, no coefficient of Excitation-S4D-Pinn rests on a box constraint in any of its six accepted identifications; its last accepted set is $B = 12.890/13.746$ inside $[4, 20]$, $C = 1.240/1.248$ inside $[1.2, 2.2]$, $D = 0.768$ on both axles inside $[0.4, 2.0]$ and $E = -2.517/-1.836$ inside $[-3, 1]$. Baseline-NN-MSE puts D on its upper bound of 2.000 on both axles in *every* accepted identification and E_f on its lower bound of -3.0 likewise. A peak factor of 2.000 against a plant realising 0.70 is less a 186% parameter error than the box-edge signature of Section III-C: the coefficient is not identified,

TABLE V
CROSS-RUN COMPARISON OF THE TWO CONFIGURATIONS.

Metric	Excitation-S4D-Pinn	Baseline-NN-MSE
<i>Closed-loop tracking (3 timed laps)</i>		
Run duration [s]	580.4	554.6
RMS e_y , overall [m]	0.256	0.319
RMS e_y , MPC [m]	0.151	0.224
ITAE e_y [ms ²]	16041	25502
MPC solve, RMS / max [ms]	0.399 / 2.604	0.380 / 2.568
<i>Axle lateral force against per-wheel telemetry</i>		
Front / rear RMSE [N]	27.2 / 33.2	329.6 / 348.4
Front / rear R^2	0.993 / 0.987	-0.160 / -0.461
Front / rear max abs. error [N]	264 / 135	1295 / 1316
<i>Peak friction against the plant’s $\mu = 0.70$</i>		
Front / rear μ RMSE, whole run	0.145 / 0.142	1.300 / 1.300
Front / rear μ RMSE, after release	0.078	1.300
<i>One-step state prediction</i>		
v_y / ω RMSE [m/s, rad/s]	0.014 / 0.049	0.156 / 0.104
Skill vs. persistence, v_y / ω	-0.27 / -8.8	-8.18 / -13.9

and no per-coefficient bounds check distinguishes that from a fit.

Second, the linear range is identified rather than inherited. Every slope is (5) for an identified set; for the plant it is the slope fitted to the simulator’s own per-wheel forces at the static axle loads $F_{z,f} = 1373$ N, $F_{z,r} = 1274$ N, which realises 17.4 kN/rad front and 16.6 kN/rad rear. The prior carries 13.2 and 13.9, 24% and 16% low; Excitation-S4D-Pinn returns 16.9 and 16.8, 3.2% low front and 1.4% high rear, removing 87% and 91% of the prior’s error. The product $BCDF_z$, the only quantity the rollout fit determines (Section V-D1), is what moves and then holds: the five cycles after the friction gate releases sit within 16.0–16.9 kN/rad, a 5% spread about a plant at 17.4. Baseline-NN-MSE returns 37.9 and 38.4, 118% and 132% high, which is $D = 2.0$ multiplying a B the fit chose against it.

Third, both sets keep the understeer sign and pass the stability gate of Section VII-A; what separates them is the grip ceiling, which (6) puts at 2.00 g for Baseline-NN-




Fig. 4. Identified tire models against the plant. Top and middle: the plant's own curve (8) at the measured static load of each axle (black) against each run's last accepted identification. Bottom: the identified Magic Formula coefficients. Baseline-NN-MSE's accepted peak factor sits on its box constraint $D = 2.0$ on both axles in every cycle, so its curve carries roughly 2.8 times the plant's force at the edge of the plotted range, while Excitation-S4D-Pinn tracks both the plant's initial slope and its peak. All force metrics are axle sums.

MSE against 0.77 g for Excitation-S4D-Pinn, a plant sustaining 0.71 g and a trajectory demanding 0.51 g. The baseline's ceiling is a claim the plant cannot honour by a factor of 2.8; only the identified set is admissible on the quantity the controller consumes. All three readings are possible because the lap now traverses part of the curvature: $P_{99}(|\alpha|) = 2.7^\circ$ and $\max |\alpha| = 5.8^\circ$ against the 8.0° at which the plant's curve peaks, at an axle utilisation of 0.71 front and 0.72 rear. This is still not a limit manoeuvre, but it suffices for the friction gate to release, which on the $\mu = 1.05$ runs it did not.

E. Peak friction

Table V reports the error of the identified peak factor against the effective friction the plant reports per wheel, and Fig. 5 its evolution over the run. This is where the pinned- D design of Section V-E2 is directly observable, because the arm starts 44 % away from the answer. Excitation-S4D-Pinn publishes six accepted identifications. The first




Fig. 5. Peak friction over the run, front axle, one row per configuration: the plant's effective μ (black) against the identified peak factor D published at each accepted identification. The rear axle is identical and is omitted, since D is pinned to the same value on both axles (Section V-E2a). The surface is static, so every departure from the black line is estimator error. Excitation-S4D-Pinn publishes the $\mu = 1.05$ prior while the friction gate is closed, then steps down onto the plant within 11 % at the first release and holds there for the remaining five cycles, while Baseline-NN-MSE sits on the $D \leq 2.0$ box constraint for the whole run.

carries $D_f = 1.009$, $D_r = 1.002$: the friction gate rejected that buffer, so the axles held the prior they came in with, the hold-rather-than-fit rule of Section V-E2a and not a fitted value. From $t = 140.6$ s the gate releases on every subsequent cycle and the published peak factor drops to 0.779, 0.821, 0.749, 0.756 and 0.768, front and rear alike, against a plant at 0.700: 10.7 % high on average, from a prior that was 44 % high. The whole-run μ RMSE is dominated by the first 140 s at the prior; over the 440 s after the first release it is 0.078 on both axles. Baseline-NN-MSE returns $D = 2.000$ on both axles in every accepted identification and never leaves the bound, a constant μ error of 1.300, so Excitation-S4D-Pinn is 9.0 $\times$ better on the whole-run figure and 16.7 $\times$ better once its warm start has released, a difference worth 1.23 g of claimed grip against a plant that sustains 0.71. Two qualifications belong on that column: the peak factor is *measured* by the brush estimator on the recorded trace and then held, not recovered by the Pacejka fit, which Section V-D1 shows cannot recover it; and the residual 10.7 % bias is one-signed, consistent with the brush model's tendency to over-read μ where the curve is not fully traversed.

F. Axle lateral force and one-step state prediction

The identified curve is scored at the measured (α, F_z) against the plant's own lateral force, summed per axle (Table V; Fig. 6(a)). Excitation-S4D-Pinn leads on both axles by an order of magnitude in RMSE and in the error tails; the baseline's negative R^2 means its curve predicts the plant's axle force worse than the mean of that force does. This is the linear-range result seen as force: at $P_{99}(|\alpha|) = 2.7^\circ$ the metric

is dominated by cornering stiffness, so the models separate almost entirely on slope. Because D multiplies that slope in (5), Baseline-NN-MSE's railed D is the direct cause of its force error. The comparison is also why the relaxed update of (14) belongs in the method rather than a configuration file: pinning D alone does not produce these numbers, since with $\beta = 1$ the same arm degrades monotonically with iteration count (Section VI-I).

Level 1 of the validation protocol is reported against a persistence baseline (next state equals last measured state) written by the same benchmark node. Excitation-S4D-Pinn leads on both signals, by an order of magnitude on lateral velocity (R^2 0.998 against 0.709) and by a factor of two on yaw rate (R^2 0.894 against 0.514); in Fig. 6(b)–(c) the separation on v_y is in the centre of the distribution, while on ω the two share a centre and differ in tail mass. Neither model beats persistence at the 0.035 s model step, the skill scores being negative for both configurations and both signals, because a one-step metric on a smooth trajectory rewards not moving. One-step accuracy is therefore weak evidence for either tire model, and the gap is read only as confirming the direction the axle-force and closed-loop comparisons establish.

G. Closed-loop tracking

Both runs finish three timed laps with zero emergency halts and the same finite-state sequence, with the handover at $t = 79.1$ s for Excitation-S4D-Pinn (Fig. 7) and $t = 64.0$ s for Baseline-NN-MSE. Under MPC, where the identified model enters the loop, Excitation-S4D-Pinn tracks at 0.151 m RMS lateral error against 0.224 m, a 32 % reduction, with maximum 0.683 m against 0.776 m (Fig. 8). ITAE is $\int t|e_y| dt$, neither scale-free nor comparable across runs of different length, so the duration-normalised forms $IAE/T = 0.104$ against 0.157 and $ITAE/(T^2/2) = 0.095$ against 0.166 are reported beside it; all three order the runs alike. The ITAE split confirms that the whole difference sits in the MPC phase: the pure-pursuit phases, which do not use the identified model, give 110 and 107, within 3 % of each other and two orders of magnitude smaller. The QP cost is unchanged: RMS and worst-case solve times are two orders of magnitude below the 50 ms control period (Fig. 7) and both slightly in the baseline's favour, by 5 % and 1 %. The identified coefficients change the model the QP carries, not its size. Lap time is not read as a result: timing runs on the ROS wall clock while the simulator advances its fixed step faster than real time, at a ratio that differed between the runs, so the 5–6 % lap-time spread is not a vehicle-performance difference.

H. Peak-friction warm start on a synthetic plant

The friction warm start is evaluated on a synthetic single-track plant with brush tires, because only a synthetic plant supplies a μ known exactly and an excitation level that can be swept independently of it: a swept-sine steering input for 40 s at 50 Hz with true $\mu = 0.70$, the same effective peak as the CARLA surface, the steering amplitude setting how far up the tire curve the manoeuvre reaches (Table VI).

TABLE VI
PEAK FRICTION RECOVERED FROM A SYNTHETIC BRUSH-TIRE PLANT
(TRUE $\mu = 0.70$), AGAINST EXCITATION.

δ amplitude	Grip utilisation	Utilisation	Brush fit (16)
0.033 rad	45 %	0.312	gated off
0.047 rad	63 %	0.442	0.699
0.060 rad	80 %	0.557	0.699
0.090 rad	100 %	0.700	0.697

The utilisation estimator tracks the grip the manoeuvre *used*, while the brush fit of (16) recovers the peak from the curvature of $F_y(\alpha)$ and is therefore flat in excitation.

The first row is refused by the absolute gate $\max |F_y|/F_z \geq 0.40$ of Section V-E2: at $\mu = 0.70$, 45 % utilisation is only 0.31 of F_z . The utilisation bound is reported instead.

At 63 % utilisation, a level an ordinary lap reaches, the utilisation bound is 37 % low while the brush fit is within 0.2 %, and it stays within 0.2 % across true $\mu \in \{0.7, 1.05, 1.4\}$ on noise-free data. With odometry and IMU noise injected the error is grip-dependent, the noise being absolute while the available force is not: at $\mu = 0.70$ it is 20 % at 63 % utilisation and 0.4 % at the peak. Below the absolute 0.40 floor on $|F_y|/F_z$ the gate suppresses the fit and reports the bound, so the failure is visible rather than silent. The gate behaves the same way on the CARLA run: it rejects the first buffer, and on the five it accepts returns μ within 11 % of the plant at an absolute axle utilisation of 0.54–0.68. That 11 % bias against the 0.2 % of the synthetic plant is the cost of a real rollout, a reconstructed $I_z \dot{\omega}$ term and a curve traversed only to 0.71 of F_z , and it dominates the identified peak factor's error.

I. What carries the result

The two arms differ in five settings at once, so that comparison establishes an ordering and attributes nothing. This subsection attributes it, at the cost of leaving the loop: every number below is an *offline replay* of the identification stage on the five 30 s buffers each arm collected during its own run, scored with the quantity the closed-loop tire-force benchmark itself reports, so these numbers and Table V's are on the same scale.

1) *Damping the co-identification loop*: Table VII sweeps β of (14) at six iterations throughout; two columns are the argument. At $\beta = 1.00$, the inherited undamped update, axle force RMSE is 70.4 and 62.3 N, reached by degrading monotonically from 38.9 and 26.3 N at a single iteration: the loop amplifies its own extrapolation instead of refining the estimate. At $\beta = 0.00$ the coefficients never move from the pinned prior, which is no identification at all, and that is the worst column, 92.5 and 71.8 N. The minimum sits between them at $\beta = 0.20$, 28.0 and 19.4 N, a mean of 23.7 N, and the identified BCD rises monotonically with β , so the damping trades how much of each fit is accepted against how much extrapolation error comes with it. Six iterations are therefore defensible on this evidence, provided the update is damped.

2) *The other three settings*: Pinning the peak factor accounts for most of what remains. Replayed on five buffers of a $\mu = 1.05$ run, letting the fit move D gives a mean $|D - \mu|$ of 0.585 and a front-stiffness spread of 32.9 kN/rad; pinning




Fig. 6. Error distributions, truth minus estimate. (a) Axle lateral force, shared bin edges over the pooled central 99% and normalised to each run's sample count; both are centred, and Baseline-NN-MSE's distribution is roughly an order of magnitude wider at every quantile (maximum absolute error 1295 N against 264 N front, 1316 N against 135 N rear). (b) One-step lateral velocity, where Excitation-S4D-Pinn is an order of magnitude tighter (0.014 m/s against 0.156 m/s RMSE). (c) One-step yaw rate, where the configurations differ in tail mass rather than centre. Neither model beats its own persistence baseline at the 0.035 s model step.




Fig. 7. Speed and solver cost along the Excitation-S4D-Pinn run: longitudinal speed v_x (left axis) against MPC QP solve time (right axis) over the full 580.4 s, shaded by the active controller state. The handover from pure pursuit to MPC falls at $t = 79.1$ s; after it the vehicle cycles the raceline's 9.6–25.8 m/s reference profile while the solve time stays two orders of magnitude below the 50 ms control period. Excitation-S4D-Pinn only; not a comparison between arms.




Fig. 8. Absolute lateral tracking error by configuration and active controller. The pure-pursuit boxes cover the acquisition phase before handover and do not use the identified model. The MPC boxes do.

only the axles whose brush fit was released reduces those to 0.191 and 32.3, and also holding the axles with no released measurement, the shipped rule, to 0.084 and 18.5 kN/rad. What removes most of the error is refusing to let an unmeasured D be fitted at all. Normalising the physics-informed objective is necessary but not sufficient alone: it moves the loss shares from 99.9996% on the dynamics residual to 1.47%, yet applied without the damping it makes the force metric slightly worse (77.1 and 76.2 N against 70.4 and 62.3 N), because an objective that now fits the data supplies a sharper

TABLE VII
RELAXATION SWEEP. FIVE RECORDED BUFFERS, SIX ITERATIONS, OFFLINE REPLAY.

β	0.00	0.05	0.10	0.15	0.20	0.25	0.30	0.50	1.00
Front F_y RMSE [N]	92.5	68.3	48.6	33.9	28.0	31.2	36.4	55.6	70.4
Rear F_y RMSE [N]	71.8	49.0	31.3	19.8	19.4	26.0	30.2	38.4	62.3
Front BCD [rad]	7.82	8.82	9.66	10.36	10.88	11.32	11.65	12.75	12.42

$\beta = 1.00$ is the inherited undamped update and $\beta = 0.00$ the pinned prior, i.e. no identification. 0.20 ships.

trajectory for the undamped loop to amplify. Two settings deserve a weaker statement than the closed-loop comparison alone would support: the S4D residual is worth 21% on the force metric at matched settings and a single iteration (35.3 N against the baseline MLP's 44.7 N), an offline replay and not a closed-loop result; and the residual *objective* is a wash, physics_informed and plain mse being indistinguishable at 23.7 N each once the normalisation and the damping are both in place. We keep it because it costs nothing and is the paper's stated formulation, not because it buys accuracy.

J. Identification cost

One S4D identification, six co-identification iterations on a 30 s lap, takes 20.9 s end to end, inside the 30 s re-identification period. The reduction from the 181.4 s it originally took is exact: no epoch count, iteration count, early-stopping rule or loss weight changed, and the identified coefficients agree to four decimal places. The per-phase breakdown is in the supplementary material.

VII. INTEGRATION WITH THE CONTROLLER

The identified model leaves the pipeline through one interface. A nonlinear MPC path-tracking controller (acados [29], partial-condensing HPIPM [30], horizon $N = 20$ at a 20 Hz control period) carries the same single-track model as the identifier and consumes Φ_p directly through (4), pushed over a ROS 2 service at the end of each identification cycle. We claim no contribution to the controller design; it appears here as the consumer that defines what an admissible identification is.

A. Admissibility gates

A coefficient set can lie strictly inside every box constraint and still be unusable. We observed two ways, each requiring a gate on a *derived* quantity rather than a coefficient:

- 1) **Grip ceiling versus trajectory demand.** The controller computes $\max_k(v_{x,k}^2|\kappa_k|)$ over the speed-clamped reference and rejects any identification whose $a_y^{\max}$ from (6) falls below it. On the earlier trajectory the degenerate fit reported a 0.40 g ceiling against a 1.22 g demand and produced 41.89 m of cross-track error at 27 m/s, against 0.03 m for a physically plausible set ($D = 1.5$) on the same trajectory, at the same speed and through the same controller; the same identified set gives 0.16 m at 11.1 m/s and 0.43 m at 15 m/s, the v^2 scaling the Introduction describes. These figures come from an offline closed-loop harness rather than a controlled pair of simulator runs, and are quoted as the observation that motivated the gate, not as a result. The gate is one-sided by construction and catches only a ceiling that is too low, so Baseline-NN-MSE's 2.00 g (Section VI-D) passes it.
- 2) **Understeer sign and critical speed.** A set with $l_f C_f > l_r C_r$ whose v_{crit} falls inside the controller's speed range makes the prediction model open-loop unstable, and the resulting $\rho(A_d) > 1$ is raised to ρ^N in the condensed QP. An observed set gave $C_f = 37031 \text{ N/rad}$ against $C_r = 11691 \text{ N/rad}$, i.e. $v_{\text{crit}} = 14.5 \text{ m/s}$ at the measured geometry and mass of Table I, inside the controller's speed range on every trajectory in this paper.

A third observation is a controller-implementation detail rather than a gate on the identification: an admissible set can still be stiff relative to the control period, so the prediction model sizes its RK4 sub-step from its own fastest lateral mode, restoring $\rho \approx 1.000$ across the speed range, at a median QP setup-and-solve cost of 0.36 ms against the 50 ms control period.

An identification that passes the gates is a model the controller can be held responsible for; one that fails is reported as unidentified rather than deployed. Two further controller-side changes were required before any identified model could be evaluated at full-scale speeds: reference speeds are clamped at ingest, and the solved state is rolled forward by the measured transport delay before each solve. Both are recorded for completeness and are not identification results.

VIII. DISCUSSION AND LIMITATIONS

A. What generalises, and what does not

The reachable-friction bound (11) is kinematic: it involves no tire parameter, no simulator-specific quantity and no property of the residual model, so it applies to any implementation of the virtual-steady-state-data strategy of [6], in simulation or on hardware. Its practical content is a design rule: the rollout must be configured to demand at least the friction one intends to identify, and where it cannot, D should be reported as unidentified rather than fitted. The diagnostic signature, coefficients resting exactly on box constraints, is likewise generic and should be checked and logged by any bounded

tire-fitting routine, being otherwise indistinguishable in the output from a successful fit. The two structural properties of Section V-D generalise on the same grounds: the BCD degeneracy follows from the quasi-steady force recovery (7) defining the virtual-data step, and the absence of contraction follows from stage 4 alone, the remedy reducing to the inherited behaviour at $\beta = 1$. What does not generalise is the value 0.20, a property of this vehicle's extrapolation error; what does is that $\beta = 1$ is not defensible and $\beta = 0$ worse still. The insufficiency of per-coefficient box constraints is also platform-independent: both failures we encountered, a set that satisfies every box constraint while degenerate and one that satisfies every box constraint while pairwise oversteering and unstable inside the controller's operating range, are properties of the parameterisation rather than of CARLA.

Three results are properties of this plant and must not be read as statements about tires: the numerical tire values, since the plant is a PhysX wheeled-vehicle model and agreement with Table II shows only that the identifier recovers the plant it observes; the gap between configured wheel friction and realised peak axle friction (1.0 against 0.70), an artefact of the server multiplying the configured coefficient by the road surface's own, which we report because an identifier tuned to converge on the configuration file's number will converge on the wrong one; and the 0.76 achieved/commanded steering ratio, where the lesson (identify against the achieved actuator state) is general but the magnitude is not. Sagmeister *et al.* [26] show sensitivity to model fidelity is largest near the acceleration limits, where this study operates, and R-CARLA [25] exists because CARLA's native dynamics are a known weak point for racing-grade work, so hardware replication is necessary before any claim about identification accuracy on a real vehicle.

B. Limitations

- 1) **One run per configuration, so no metric in Table V carries an uncertainty.** Every value is a point estimate from a single 580.4 s and a single 554.6 s run. The 0.151 m against 0.224 m RMS tracking difference in particular has no confidence interval and no evidence that it exceeds run-to-run variation, which is not measured here; the order-of-magnitude force and friction margins are large enough that such variation is unlikely to reverse them, but that is an argument, not a measurement.
- 2) **Single surface, vehicle and trajectory.** That the surface differs from the one the prior was measured on makes the identification a real result rather than a restatement of the prior, but it is *not* a demonstration of adaptivity: the identifier is never asked to follow a *change* in friction. The runtime friction override and the step and decay schedules the scenario file already carries support that experiment, and we do not perform it here.
- 3) **The friction estimator's excitation sweep is synthetic** (Table VI). On the CARLA runs the warm start is exercised at the one excitation level the lap supplies, returning μ 11 % high against the 0.2 % of the synthetic plant, so the excitation dependence of that bias is not

mapped against the simulator's per-wheel forces. The published peak factor is exactly this estimate, so the bias passes through undamped.

- 4) **The rear axle is not identified on this trajectory** and the front-axle μ is applied to both (Section V-E2a). The quasi-steady argument making $\mu_f = \mu_r$ holds on a homogeneous surface and would not survive a split- μ one, which this sensor set cannot detect in the first place.
- 5) **The two configurations differ in five settings at once** (Table III), so Section VI-C attributes nothing to the residual architecture, the loss, the solver, the warm start or the relaxation individually. Section VI-I does attribute them, but by offline replay on recorded buffers, so its ordering is not a closed-loop ordering. The S4D figure there (21 % on force RMSE at matched settings) has the same status: we report it as a measured ablation and not as a reproduction of the closed-loop S4 improvement of [20]. The excitation-aware rollout, this paper's central contribution, is active in *both* arms, so the closed-loop comparison is not evidence for the identifiability correction; the evidence for that is the sweep of Section VI-B.
- 6) **The physics-informed objective buys nothing measurable here.** Once its units are reconciled with the data term and the loop is damped, it is indistinguishable from a plain one-step MSE on this manoeuvre. It is retained because it is the paper's stated formulation and costs nothing, not because it is shown to help.
- 7) **Longitudinal dynamics are out of scope.** As in the baseline, only lateral dynamics are identified. Combined slip and load transfer are neglected, and the simulator's longitudinal wheel force channel is drivetrain torque rather than a contact-patch force, so it is unusable as ground truth for a longitudinal extension without further work.

IX. CONCLUSION AND FUTURE WORK

This paper studied the identifiability of a learning-based on-track Pacejka identification pipeline [6] at full vehicle scale, using a CARLA simulation architecture that exposes the plant's own per-wheel forces as ground truth. Transplanted unchanged from a 1:10 platform, the pipeline fails in a specific, diagnosable and correctable way: a loss of excitation created by its own virtual-data step, which no change of solver, bounds or tuning removes, and whose reachable friction depends on the rollout configuration and not the tire. Two further properties of the loop follow from the same analysis. The fit's residual depends on B , C and D only through their product, because stage 3 reconstructs its forces from the rollout's own quasi-steady yaw balance and so observes nothing but the lateral acceleration the previous coefficient set produced; the peak factor is therefore not recoverable from the rollout at any excitation and must be measured on the recorded trace and held. And because each fit is handed the trajectory its own previous answer generated, the loop is a positive feedback path rather than a contraction, so stage 4 is reformulated as a relaxed fixed-point iteration.

The correction comes with an extrapolation bound on the steering sweep, a frozen regularisation target, identification against the achieved rather than the commanded road-wheel angle, measured mass and geometry, a friction warm start identified from the curvature of the axle force curve and pinned for the loop, a relaxation $\beta = 0.20$, and a dynamics residual normalised to the units of the data term. Together these produce identifications with no railed coefficient and the correct understeer sign, at the margins Section VI reports. Those numbers are an identification and not a preserved prior, with three qualifications: the peak factor is identified by the brush estimator on the recorded trace, not by the Pacejka fit, which cannot identify it (Section V-D1); one surface is not adaptivity, which needs a friction *change* the runs do not contain; and each configuration is a single run, so the closed-loop margins are observations rather than estimates with a confidence interval. We further argue that per-coefficient box constraints are not a sufficient admissibility criterion for handing a model to a model-based controller, and place gates on derived physical quantities at that interface instead.

A. Future work

- 1) **Sweep the friction rather than merely mismatching it.** These runs put the prior 44 % away from a *static* surface, showing the pipeline moves on evidence but not that it tracks; the step and decay schedules already in the scenario file are the intended vehicle, and the strongest available argument for identifying on track at all.
- 2) **Run the ablations of Section VI-I in the loop.** The relaxation sweep, the D -pinning comparison and the S4D and objective ablations are offline replays, enough to attribute the result but not to confirm the ordering survives closed-loop coupling. Each is a scenario-file row.
- 3) **Excite the peak, and design for it.** The lap reaches $P_{99}(|\alpha|) = 2.7^\circ$ against a plant whose curve peaks at 8.0° ; that is enough for the friction gate to release and is also where the brush estimator's 11 % bias comes from. Equation (11) inverts into a requirement on the manoeuvre, so a controller that deliberately inserts bounded slalom or ramp-steer excitation when the identifier reports insufficient slip-angle coverage would move the pipeline toward active, safety-bounded experiment design [49].
- 4) **Hardware replication.** The decisive test of every claim in Section VIII is a physical Formula Student vehicle. The identification code is unchanged between simulation and hardware; what changes is the availability of ground truth, and that is what makes the simulation study a necessary precursor rather than a substitute.

APPENDIX A

ASSUMPTIONS AND REPRODUCIBILITY NOTES

The following are assumptions or design choices rather than measured facts, listed so a reader can judge how far each result depends on them.

A1. *Yaw inertia.* $I_z = 51.1 \text{ kg m}^2$ is carried in the model file and is *not* measured. It enters (2) directly, so it scales the yaw-rate residual the network is trained on. It can be discharged by applying a known yaw moment and fitting I_z from the measured $\dot{\omega}$, or by reading the actor's inertia tensor.

A2. *Centre-of-gravity split.* $l_f = 0.738142 \text{ m}$, $l_r = 0.795362 \text{ m}$ are taken from the controller configuration. The measured static load split (51.2/48.8) implies a marginally different split than the geometry does (51.9/48.1), a 1.4 % disagreement, below the level at which any conclusion here would change.

A3–A4. *CG height and steering sign.* $h_{cg} = 0.3 \text{ m}$ is carried but unused: the single-track model neglects load transfer, following the baseline, and normal loads are static. `feedback/steering_angle` is published in degrees and treated as the front road-wheel angle, positive left, verified against the front-left wheel joint state to within 0.5 %; left and right road wheels are not distinguished.

A5. *Steady-state force recovery.* Equation (7) is not a force measurement; being fully determined by $v_x \omega$ it misattributes transients, and it is applied only to the *synthetic*, quasi-steady rollout. Its adequacy was checked on measured data against the simulator's own axle force (correlation 0.993, slope 0.94).

A6. *The residual network is not the plant.* The corrected model is queried by the rollout at inputs the network has only partially seen. The extrapolation bound of Section V-C limits but does not eliminate this: the sweep still reaches $1.5 \times$ the observed 99th percentile of $|\delta|$, a choice that is empirical rather than derived.

A7. *Simulator determinism and timing.* CARLA runs in synchronous mode with a 0.033 s physics step; sim time advances faster than wall time ($\approx 1.9 \times$ on our hardware), and all identification data is stamped from the simulation clock.

A8. *Software configuration of record.* ROS 2 Humble, Python 3.10, PyTorch with CUDA on an RTX 2080-class GPU; the bridge is a separate overlay; CARLA server 0.9.16 / Unreal Engine 4.26; acados 0.5.5 with partial-condensing HPIPM. The identifier's acquisition timer runs at 50 Hz (20 Hz is the controller's period) with a 30 s buffer, six co-identification iterations, 200 training epochs, full-batch Adam at 5×10^{-4} , and $T_s = 0.035 \text{ s}$. `nn_train` seeds PyTorch and NumPy at 42; `pacejka_solver.seed` is null in both shipped parameter sets, so the reported identifications are reproducible in distribution but not bit for bit, while the sweep of Table IV is seeded explicitly and is bit-reproducible.

APPENDIX B

SIMULATOR–ROS 2 INTERFACE

Feedback path

Table VIII lists the interface consumed by the identifier and its validation.

a) *Per-wheel ground truth:* `feedback/tire_forces` carries CARLA's own `WheelTelemetryData` in wheel order [FL, FR, RL, RR]. No analytic tire model exists inside the bridge, so nothing on this topic can disagree with the physics the vehicle is driven by. Four measured properties bound what may be

TABLE VIII
ROS 2 INTERFACE USED BY THE IDENTIFICATION AND VALIDATION PIPELINE. NAMESPACE IS `/sim` UNLESS MARKED ABSOLUTE.

Topic	Type	Role
<code>/odom</code>	<code>nav_msgs/Odometry</code>	State v_x, v_y, ω ; twist is body-frame
<code>feedback/steering_angle</code>	<code>std_msgs/Float32 (deg)</code>	Achieved road-wheel angle δ
<code>feedback/imu</code>	<code>sensor_msgs/Imu</code>	a_x, a_y for the friction warm-start
<code>feedback/tire_forces</code>	<code>sim_manager_msgs/TireForces</code>	Per-wheel $\alpha, F_z, F_y, \mu, \omega_w, \kappa$: ground truth
<code>feedback/vehicle_physics</code>	<code>std_msgs/String (JSON)</code>	Live wheel parameters of (8)
<code>feedback/speed</code>	<code>std_msgs/Float32</code>	Forward speed echo
<code>feedback/motors</code>	<code>std_msgs/String (JSON)</code>	Per-wheel speed/torque/brake
<code>/drive</code>	<code>ackermann_msgs/AckermannDriveStamped</code>	Control command in
<code>control/tire_friction</code>	<code>std_msgs/Float32</code>	Runtime grip override
<code>/clock (abs.)</code>	<code>roscpp_msgs/Clock</code>	Simulation time base

read off it. (i) Only `lateral_force` is force ground truth; it is the channel every validation in this paper uses. (ii) `longitudinal_force` is a friction-capacity report rather than a measurement, tracking μF_z and correlating with ma_x at 0.04, and `wheel_torque` is an exact identity of it. (iii) Below 0.5 m/s the publisher zeroes slip angle and both forces while the normal load stays live, so the identifier gates on the all-zero pattern as well. (iv) The slip angle needs no sign change across the boundary, verified rather than assumed.

Figure 9 states the conventions the rest of the paper assumes. CARLA inherits Unreal's left-handed world while every ROS 2 topic follows REP-103 [50] and the identification equations are written in the body frame with α and F_y positive to the left, the SAE-style axis convention [51], [4]. The bridge is a per-channel conversion rather than a single transform, and the trap is that the tire telemetry's slip angle needs *no* sign change, because the frame flip and CARLA's own convention cancel: an error there is invisible in the states and inverts the identified curve, so it was checked against a slip angle rebuilt from the chassis state in a steady 4 m/s turn, equal in magnitude to within 2 % front and 15 % rear.

APPENDIX C

REFERENCE IMPLEMENTATION

The identification package, the ROS 2 bridge and the controller stack used in this paper are available as `On-Track-SysID`, extended from [6] (upstream: <https://github.com/ForzaETH/On-Track-SysID>), `CARLA_ASU_BRIDGE` (https://github.com/ebrahimabdelghfar/Carla_ASU_Bridge) [37], and the accompanying MPC and benchmarking packages, released together with the scenario definitions, the exported result CSVs and the figure and sweep generators at https://github.com/ebrahimabdelghfar/Ebrahim_Master_Thesis_repo/tree/feat-follow-premade-racing-track.




Fig. 9. Frame and sign conventions across the simulator-ROS 2 boundary. (a) All identification quantities are defined in the body frame: x forward, y left, ω and δ positive counter-clockwise, and slip angle and lateral force positive to the left, so that a positive α produces a positive F_y [51]. (b) The conversion applied per channel. The slip angle of the tire telemetry is the one exception that needs no sign change, because the frame flip and CARLA's own convention cancel.

ACKNOWLEDGMENT

The ON-TRACK-SYSID pipeline studied here originates from the ForzaETH Race Stack and the work of Dikici *et al.* [6], released under an open-source licence; this paper reports a port, a correction, and an extension of that work rather than a reimplementaion from scratch.

FUNDING

This work received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

CONFLICT OF INTEREST

The authors declare no conflict of interest.

ETHICS

All results in this paper are obtained in simulation. No experiment involved human subjects, personal data, or the operation of a physical vehicle, so no human-subjects approval and no vehicle-safety protocol were required.

DATA AND CODE AVAILABILITY

The identification package, the ROS 2 bridge, the MPC and the benchmarking packages, the scenario definitions, the exported result CSVs behind every figure and table of Section VI, and the sweep generator behind Table IV are in the repositories listed in Appendix C; the top-level project repository is https://github.com/ebrahimabdelghfar/Ebrahim_Master_Thesis_repo/tree/feat-follow-premade-racing-track. ON-TRACK-SYSID is an extension of the ForzaETH release of [6] and is distributed under the MIT licence it declares in its package manifest; CARLA_ASU_BRIDGE is GPL-3.0. The two are used as separate ROS 2 overlays communicating only over ROS 2 topics and services, and neither links against the other, so no combined-work obligation arises between them; the identification and control packages in this work are released under the MIT terms the ON-TRACK-SYSID upstream requires.

REFERENCES

- [1] A. Liniger, A. Domahidi, and M. Morari, "Optimization-based autonomous racing of 1:43 scale RC cars," *Optimal Control Applications and Methods*, vol. 36, no. 5, pp. 628–647, 2015.
- [2] J. Kabzan, L. Hewing, A. Liniger, and M. N. Zeilinger, "Learning-based model predictive control for autonomous racing," *IEEE Robotics and Automation Letters*, vol. 4, no. 4, pp. 3363–3370, 2019.
- [3] N. Baumann, E. Ghignone, J. Kühne, N. Bastuck, J. Becker, N. Imholz, T. Kränzlin, T. Y. Lim, M. Lötscher, L. Schwarzenbach, L. Tognoni, C. Vogt, A. Carron, and M. Magno, "ForzaETH race stack—scaled autonomous head-to-head racing on fully commercial off-the-shelf hardware," *Journal of Field Robotics*, vol. 42, no. 4, pp. 1037–1079, 2024, preprint: arXiv:2403.11784.
- [4] H. B. Pacejka, *Tire and Vehicle Dynamics*, 3rd ed. Oxford, UK: Butterworth-Heinemann, 2012.
- [5] J. Kabzan, M. I. Valls, V. J. F. Reijgwart, H. F. C. Hendriks, C. Ehmke, M. Prajapat, A. Bühler, N. Gosala, M. Gupta, R. Sivanesan, A. Dhall, E. Chisari, N. Karnchanachari, S. Brits, M. Dangel, I. Sa, R. Dubé, A. Gawel, M. Pfeiffer, A. Liniger, J. Lygeros, and R. Siegwart, "AMZ driverless: The full autonomous racing system," *Journal of Field Robotics*, vol. 37, no. 7, pp. 1267–1294, 2020.
- [6] O. Dikici, E. Ghignone, C. Hu, N. Baumann, L. Xie, A. Carron, M. Magno, and M. Corno, "Learning-based on-track system identification for scaled autonomous racing in under a minute," *IEEE Robotics and Automation Letters*, vol. 10, no. 2, 2025, preprint: arXiv:2411.17508. Code: <https://github.com/ForzaETH/On-Track-SysID>.
- [7] L. Ljung, *System Identification: Theory for the User*, 2nd ed. Upper Saddle River, NJ, USA: Prentice Hall, 1999.
- [8] A. Dosovitskiy, G. Ros, F. Codevilla, A. López, and V. Koltun, "CARLA: An open urban driving simulator," in *Proc. 1st Annual Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 78. PMLR, 2017, pp. 1–16. [Online]. Available: <https://proceedings.mlr.press/v78/dosovitskiy17a.html>
- [9] E. Bakker, L. Nyborg, and H. B. Pacejka, "Tyre modelling for use in vehicle dynamics studies," SAE International, SAE Technical Paper 870421, 1987.
- [10] T. A. Wenzel, K. J. Burnham, M. V. Blundell, and R. A. Williams, "Dual extended Kalman filter for vehicle state and parameter estimation," *Vehicle System Dynamics*, vol. 44, no. 2, pp. 153–171, 2006.
- [11] C. Sierra, E. Tseng, A. Jain, and H. Peng, "Cornering stiffness estimation based on vehicle lateral dynamics," *Vehicle System Dynamics*, vol. 44, no. suppl, pp. 24–38, 2006.
- [12] R. Rajamani, D. Piyabongkarn, J. Y. Lew, K. Yi, and G. Phanomchoeng, "Algorithms for real-time estimation of individual wheel tire-road friction coefficients," *IEEE/ASME Transactions on Mechatronics*, vol. 17, no. 6, pp. 1183–1195, 2012.
- [13] C. Ahn, H. Peng, and H. E. Tseng, "Robust estimation of road frictional coefficient," *IEEE Transactions on Control Systems Technology*, vol. 21, no. 1, pp. 1–13, 2013.

- [14] M. Acosta, S. Kanarachos, and M. Blundell, "Road friction virtual sensing: A review of estimation techniques with emphasis on low excitation approaches," *Applied Sciences*, vol. 7, no. 12, p. 1230, 2017.
- [15] S. Khaleghian, A. Emami, and S. Taheri, "A technical survey on tire-road friction estimation," *Friction*, vol. 5, no. 2, pp. 123–146, 2017.
- [16] Y.-H. J. Hsu, S. M. Laws, and J. C. Gerdes, "Estimation of tire slip angle and friction limits using steering torque," *IEEE Transactions on Control Systems Technology*, vol. 18, no. 4, pp. 896–907, 2010.
- [17] L. Hewing, A. Liniger, and M. N. Zeilinger, "Cautious NMPC with Gaussian process dynamics for autonomous miniature race cars," in *Proc. European Control Conference (ECC)*, 2018, pp. 1341–1348.
- [18] L. Hewing, J. Kabzan, and M. N. Zeilinger, "Cautious model predictive control using Gaussian process regression," *IEEE Transactions on Control Systems Technology*, vol. 28, no. 6, pp. 2736–2743, 2020.
- [19] J. Chrosniak, J. Ning, and M. Behl, "Deep dynamics: Vehicle dynamics modeling with a physics-constrained neural network for autonomous racing," *IEEE Robotics and Automation Letters*, vol. 9, no. 6, pp. 5292–5297, 2024, preprint: arXiv:2312.04374.
- [20] Z. Wu, C. Hu, Y. Wang, L. Xie, and H. Su, "Vision-augmented on-track system identification for autonomous racing via attention-based priors and iterative neural correction," 2026. [Online]. Available: <https://arxiv.org/abs/2603.09399>
- [21] A. Gu, K. Goel, A. Gupta, and C. Ré, "On the parameterization and initialization of diagonal state space models," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022, preprint: arXiv:2206.11893.
- [22] A. Gu, K. Goel, and C. Ré, "Efficiently modeling long sequences with structured state spaces," in *Proc. International Conference on Learning Representations (ICLR)*, 2022, preprint: arXiv:2111.00396.
- [23] E. Fiala, "Seitenkräfte am rollenden luftreifen," *VDI-Zeitschrift*, vol. 96, no. 29, pp. 973–979, 1954.
- [24] NVIDIA Corporation, "PhysX SDK 4.1 documentation: Vehicles — tire force computation (PxVehicleComputeTireForceDefault)," Software documentation, PhysX SDK 4.1, as bundled with Unreal Engine 4.26 and linked by CARLA 0.9.16, 2019, <https://github.com/NVIDIA-Omniverse/PhysX/tree/release/104.2/physx>.
- [25] M. Brunner, E. Ghignone, N. Baumann, and M. Magno, "R-CARLA: High-fidelity sensor simulations with interchangeable dynamics for autonomous racing," 2025. [Online]. Available: <https://arxiv.org/abs/2506.09629>
- [26] S. Sagmeister, P. Kounatidis, S. Goblirsch, and M. Lienkamp, "Analyzing the impact of simulation fidelity on the evaluation of autonomous driving motion control," 2026. [Online]. Available: <https://arxiv.org/abs/2602.07984>
- [27] J. Abdo, A. Shibu, M. Saeed, A. M. Aga, A. Sivaprazad, and M. Al-Musleh, "Simulating an autonomous system in CARLA using ROS 2," 2025. [Online]. Available: <https://arxiv.org/abs/2511.11310>
- [28] M. O'Kelly, H. Zheng, D. Karthik, and R. Mangharam, "FITENTH: An open-source evaluation environment for continuous control and reinforcement learning," in *Proc. NeurIPS 2019 Competition and Demonstration Track*, ser. Proceedings of Machine Learning Research, vol. 123. PMLR, 2020, pp. 77–89. [Online]. Available: <https://proceedings.mlr.press/v123/o-kelly20a.html>
- [29] R. Verschuere, G. Frison, D. Kouzoupis, J. Frey, N. van Duijken, A. Zanelli, B. Novoselnik, T. Albin, R. Quirynen, and M. Diehl, "acados—a modular open-source framework for fast embedded optimal control," *Mathematical Programming Computation*, 2022, preprint: arXiv:1910.13753.
- [30] G. Frison and M. Diehl, "HPIPM: A high-performance quadratic programming framework for model predictive control," *IFAC-PapersOnLine (Proc. 21st IFAC World Congress)*, vol. 53, no. 2, pp. 6563–6569, 2020.
- [31] A. Heilmeyer, A. Wischniewski, L. Hermansdorfer, J. Betz, M. Lienkamp, and B. Lohmann, "Minimum curvature trajectory planning and control for an autonomous race car," *Vehicle System Dynamics*, vol. 58, no. 10, pp. 1497–1527, 2020.
- [32] R. Rajamani, *Vehicle Dynamics and Control*, 2nd ed. New York, NY, USA: Springer, 2012.
- [33] R. Bellman and K. J. Åström, "On structural identifiability," *Mathematical Biosciences*, vol. 7, no. 3, pp. 329–339, 1970.
- [34] A. Raue, C. Kreutz, T. Maiwald, J. Bachmann, M. Schilling, U. Klingmüller, and J. Timmer, "Structural and practical identifiability analysis of partially observed dynamical models by exploiting the profile likelihood," *Bioinformatics*, vol. 25, no. 15, pp. 1923–1929, 2009.
- [35] K. J. Åström and B. Wittenmark, *Adaptive Control*, 2nd ed. Reading, MA, USA: Addison-Wesley, 1995.
- [36] P. Kaur, S. Taghavi, Z. Tian, and W. Shi, "A survey on simulators for testing self-driving cars," in *Proc. 4th Int. Conf. on Connected and Autonomous Driving (MetroCAD)*, 2021, pp. 62–70. [Online]. Available: <https://arxiv.org/abs/2101.05337>
- [37] E. Abdelghfar, "Formula AI simulator: A CARLA-based simulation bridge for autonomous formula student vehicles," Software, GPL-3.0, 2026, https://github.com/ebrahimabdelghfar/Carla_ASU_Bridge.
- [38] D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," in *International Conference on Learning Representations (ICLR)*, 2015.
- [39] A. Paszke, S. Gross, F. Massa, *et al.*, "PyTorch: An imperative style, high-performance deep learning library," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [40] M. Raissi, P. Perdikaris, and G. E. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019.
- [41] J. A. Nelder and R. Mead, "A simplex method for function minimization," *The Computer Journal*, vol. 7, no. 4, pp. 308–313, 1965.
- [42] R. Storn and K. Price, "Differential evolution – a simple and efficient heuristic for global optimization over continuous spaces," *Journal of Global Optimization*, vol. 11, no. 4, pp. 341–359, 1997.
- [43] M. A. Branch, T. F. Coleman, and Y. Li, "A subspace, interior, and conjugate gradient method for large-scale bound-constrained minimization problems," *SIAM Journal on Scientific Computing*, vol. 21, no. 1, pp. 1–23, 1999.
- [44] P. Virtanen, R. Gommers, T. E. Oliphant, *et al.*, "SciPy 1.0: Fundamental algorithms for scientific computing in Python," *Nature Methods*, vol. 17, pp. 261–272, 2020.
- [45] A. Savitzky and M. J. E. Golay, "Smoothing and differentiation of data by simplified least squares procedures," *Analytical Chemistry*, vol. 36, no. 8, pp. 1627–1639, 1964.
- [46] J. Svendenius and M. Gäfvert, "A brush-model-based semi-empirical tire-model for combined slips," SAE International, SAE Technical Paper 2004-01-1064, 2004.
- [47] B. P. Welford, "Note on a method for calculating corrected sums of squares and products," *Technometrics*, vol. 4, no. 3, pp. 419–420, 1962.
- [48] Institute of Automotive Technology, Technical University of Munich, "global_racetrajectory_optimization," Software, Institute of Automotive Technology, Technical University of Munich, 2019–2022, https://github.com/TUMFTM/global_racetrajectory_optimization.
- [49] G. C. Goodwin and R. L. Payne, *Dynamic System Identification: Experiment Design and Data Analysis*, ser. Mathematics in Science and Engineering. New York, NY, USA: Academic Press, 1977, vol. 136.
- [50] T. Foote and M. Purvis, "Standard units of measure and coordinate conventions," Open Source Robotics Foundation, ROS Enhancement Proposal 103, 2010, <https://www.ros.org/reps/rep-0103.html>.
- [51] SAE International, "Vehicle dynamics terminology," SAE International, Surface Vehicle Recommended Practice J670, 2008.